

Multimodal Conversational AI for E-commerce: A Vision-Language Approach

Matteo Mirgone — MSADS, University of Chicago

2026-05-19

Multimodal Conversational AI for E-commerce: A Vision-Language Approach

Course: Generative AI, Principles & Applications — MSc Applied Data Science, University of Chicago

Author: Matteo Mirgone **Date:** 2026-05-19

Abstract

This project implements a multimodal conversational AI system for e-commerce product Q&A on the Amazon Product Dataset 2020. The system combines CLIP-based multimodal embeddings, vector retrieval through ChromaDB, and grounded response generation via the Llama 3.1 large language model. Users can query the catalog with natural-language text, an uploaded product image, or both modalities together. Because retrieval operates entirely within CLIP's shared text-image embedding space, the system serves text, image, and mixed queries from a single index with no dedicated vision-language model on the generation side. Quantitative evaluation across 500 held-out products from a 1,991-product indexed catalog yields **Recall@5 of 0.998** for realistic partial-text queries and **Recall@5 of 0.816** for cross-modal image-to-text retrieval. The end-to-end system is delivered as a Streamlit web application that handles the example interactions specified in the project brief — text-based product questions, image-based product identification, and image-retrieval requests — under zero-shot, few-shot, and multi-shot prompting strategies.

1. Introduction

E-commerce customer support has historically been constrained by single-modality interfaces. Text-based assistants cannot interpret an uploaded image of a product, and image-based systems struggle with nuanced natural-language queries about specifications or comparisons. Customers, meanwhile,

naturally express product questions in mixed-modality ways: “what is this?” alongside a photo, or “show me what the AirPods Pro look like” in plain text.

A multimodal conversational system that processes both text and images bridges this gap. By leveraging a model that aligns visual and textual representations in a shared embedding space, retrieval can serve text queries, image queries, and combinations of both with a single index. The retrieved product information then grounds a language model’s response, ensuring answers stay tied to actual catalog data rather than the LLM’s parametric memory.

This project builds such a system end-to-end, comprising four major components:

1. **Data preprocessing** of the Amazon Product Dataset 2020 to produce a clean, indexable catalog with cached primary images
2. **A Vision-Language Retrieval-Augmented Generation (RAG) framework** using CLIP (`openai/clip-vit-base-patch32`) as the multimodal encoder and ChromaDB as the persistent vector store
3. **LLM integration** with Llama 3.1 8B (via Groq), supporting zero-shot, few-shot, and multi-shot prompting strategies for context-grounded response generation
4. **A Streamlit user interface** that accepts text and image input simultaneously and displays both the generated response and the retrieved products

2. Background and related work

2.1 CLIP

CLIP (Contrastive Language–Image Pre-training; Radford et al., 2021) is a foundation model trained on 400 million image-text pairs from the web. Its architecture comprises two separate encoders — a Vision Transformer for images and a Transformer for text — trained with a contrastive objective that pulls matched image-text pairs together in a shared 512-dimensional embedding space while pushing mismatched pairs apart. The resulting embedding space exhibits a property critical for this project: text and images describing the same concept land near each other. A text query “wireless earbuds” and an image of AirPods produce embeddings with high cosine similarity, allowing both to retrieve from the same vector index.

This project uses the `openai/clip-vit-base-patch32` variant, which trades some accuracy for substantially faster inference compared to the larger `ViT-L/14` variant. For catalog retrieval at the ~2K-product scale used here, this trade-off is favorable: total encoding time was under 8 minutes on a Colab T4 GPU, and retrieval performance (see §5) is already near-saturated on the realistic-query regime.

2.2 Retrieval-Augmented Generation

RAG (Lewis et al., 2020) addresses two well-known LLM failure modes — hallucination of facts and inability to access information outside the training corpus — by injecting retrieved external context into the prompt at inference time. The standard pipeline embeds the user query, retrieves the top-k most similar documents from a vector index, formats them into the LLM’s context window, and generates a response that grounds its claims in the retrieved material.

This project applies RAG to a multimodal setting: the query may be text, an image, or both; the retrieved items are products with both textual descriptions and visual representations; and the generation model uses the textual portion of retrieved products as grounding context.

2.3 Open-source LLMs

The project specification permits “an open-source LLM (e.g., Meta-Llama-3.1 or Mixtral).” This implementation uses **Llama 3.1 8B Instruct**, accessed through Groq’s hosted inference service. The 8B variant is sufficient for grounded RAG tasks where the answer is present in the retrieved context — a setting that benefits less from a 70B model’s superior parametric knowledge than from low-latency response generation. Groq’s free tier provides sub-second response times at this prompt size, which is what makes the multi-shot strategy practical in the UI without noticeable lag.

3. Data

3.1 Source

The project uses the **Amazon Product Dataset 2020** from Kaggle (promptcloud). The raw CSV is ~19 MB and contains roughly 10K product rows spanning a long-tail of categories. Each row contains a product name, brand, category path (| -separated), pricing, free-text description (“About Product”), technical specifications, and a | -separated list of image URLs.

3.2 Preprocessing

Phase 1 of the pipeline (notebooks/build_index_colab.ipynb , sections 1–13) performs the following steps:

1. **Schema audit.** Per-column missingness, uniqueness, and sample values are computed to identify usable fields. The dataset has substantial null rates in many columns (e.g., `Brand Name` is missing for roughly half of all rows, `Selling Price` and `About Product` for ~10–20%). Required fields are reduced to those with reliable coverage.
2. **Required-field filtering.** Rows without a product name or image URL are dropped, since both modalities are required for indexing.

3. **Deduplication.** Rows with identical (Product Name, Brand Name) pairs are reduced to a single entry, removing variants listed multiple times under the same name.
4. **Description construction.** Three candidate description strings are built per product, varying in scope:
 - `desc_minimal` : title + brand
 - `desc_standard` : title + brand + category + price + about-product (*this is the variant indexed*)
 - `desc_full` : standard + technical details + specificationsCLIP's text encoder truncates at 77 tokens (~50–60 words), which limits the practical benefit of the `desc_full` variant. The mean character length of the indexed `desc_standard` field is 527 characters, of which only the first ~250–300 characters reach the text encoder after tokenization.
5. **Image URL parsing.** The `Image` field is `|`-separated; the first URL is retained as the primary image.
6. **Stratified sampling.** To keep encoding compute tractable, the cleaned catalog is reduced to a target size with proportional sampling per top-level category and a floor of 30 products per category to maintain representativeness of long-tail categories. The actual final size (after image-download failures, see step 7) is **1,991 products across 24 top-level categories**.
7. **Image caching.** Primary images are downloaded in parallel (24 workers), resized so the maximum side is 512 pixels, and stored as JPEG. **1,987 of 1,991 (99.8%) downloaded successfully**; the four failures were `requests.ReadTimeout` errors on slow CDN responses, not 404s. The cached images live at `prepared_data/images/<product_id>.jpg`.

3.3 Final dataset

After all preprocessing steps, the indexed catalog contains:

- **1,991 products** with both a description string and a cached primary image (1,987 images on disk; 4 referenced but missing)
- **24 top-level categories**, with a heavy skew toward Toys & Games (1,207 products / 60.6% of the catalog) — a real property of the source dataset that downstream evaluation has to account for
- **527-character mean description length** for the indexed `desc_standard` field

The top-five categories in the indexed catalog:


Category	Count
Toys & Games	1,207
Unknown (missing/malformed category)	148
Home & Kitchen	129
Clothing, Shoes & Jewelry	113
Sports & Outdoors	97

This category imbalance is the dominant property of the data, and it shapes every qualitative result in §5: queries that match toys/games concepts retrieve well, while queries about other categories (electronics, apparel) compete with a much smaller pool of in-domain products.

4. System design

4.1 Overview

The system comprises two pipelines: an offline indexing pipeline (Phases 1–2) and an online query pipeline (Phases 4–5). Both share the same CLIP model and the same vector database, providing a single point of truth for what the system “knows” about the catalog.



4.2 Embedding generation

CLIP is loaded once and used in two contexts: - **During indexing**, every product’s `desc_standard` is text-encoded and every product’s primary image is image-encoded. Both embeddings are L2-

normalized so that downstream cosine similarity reduces to a dot product. - **During inference**, the user's query is encoded — text via `CLIPModel.get_text_features`, image via `CLIPModel.get_image_features` — using the same model and the same normalization.

Encoding is batched (batch size 64 for text, 32 for images) under `torch.no_grad()`. End-to-end encoding of 1,991 products took ~8 minutes on a Colab T4 GPU.

Note on dependency pinning. `transformers` is pinned to `>=4.36,<5`. In `transformers 5.x`, `get_text_features` / `get_image_features` were changed to return a structured output object rather than a tensor, which silently breaks every downstream `.cpu().numpy()` call in this pipeline. 4.x is required.

4.3 Vector storage

ChromaDB is used as the persistent vector store. Two collections are created:

- `products_text` — 512-dimensional text embeddings, one per product
- `products_image` — 512-dimensional image embeddings, one per product

Both use cosine similarity for retrieval (`hnsw:space="cosine"`). Each entry's metadata includes the product ID, name, brand, category, price, about-product text, local image path, and original image URL — the fields needed to render results in the UI and to construct the LLM prompt.

The project specification names Vertex AI Vector Search as an example vector store. ChromaDB was selected here for three pragmatic reasons: zero infrastructure cost, full local persistence (the entire 1,991-product index fits in a 24 MB SQLite file), and the same HNSW cosine-similarity algorithm that backs Vertex. The retrieval interface in `app/rag_chain.py` is thin enough that swapping to Vertex AI Vector Search would require approximately 50 lines of code — the embedding generation and prompt-construction logic are unchanged.

4.4 Retrieval

The retrieval pipeline in `app/rag_chain.py::MultimodalRAG.retrieve` is hybrid: it combines CLIP vector search with a keyword-name index and merges the two via Reciprocal Rank Fusion. Three precedence rules are then applied on top of the fused list. The full ordering of decisions per query:

1. **Vector retrieval (CLIP)** — encode the text query with CLIP-text and the image (if any) with CLIP-image. Fetch top-20 candidates per modality from the corresponding ChromaDB collection. Cosine similarity is exposed as `1 - distance` (ChromaDB returns distance for cosine space).
2. **Keyword retrieval** — for text queries, run an in-memory keyword search over product names and categories. Each product is scored by how many query keywords (≥ 3 chars, non-stopword, singular/plural-aware) appear in its name. Multi-keyword matches are weighted super-linearly ($(\text{hits} / |\text{kws}|)^2$) so a 2-of-2 match dominates a 1-of-2 match.

3. **Reciprocal Rank Fusion (RRF)** — the CLIP vector ranking, keyword ranking, and image ranking are merged with the RRF scoring formula $\sum 1 / (rrf_c + rank)$, $rrf_c = 20$. The smaller-than-default constant amplifies top-rank contributions, which matters for product search where the right answer should be among the top 3, not the top 60.
4. **Lexical anchor override** — after RRF, if the user typed a query with ≥ 2 meaningful keywords and any retrieved product contains *all* of those keywords as whole-word tokens in its name, that product is promoted to position 0. This catches the case where CLIP semantically prefers an ambiguous match (e.g. “Cat Accessory Kit” — a girl in a cat costume) over the literal lexical match (“Amscan Miss Meow Cat Costume”).
5. **Strong-image-match override** — if an image was uploaded and its CLIP-image top hit has cosine ≥ 0.95 (the `EXACT_MATCH` threshold), that hit is promoted to position 0 unconditionally. This is the regime where the user has uploaded a photo of a product that is in the catalog: the system must identify it correctly even if the accompanying text query is vague (e.g. “*price of*”).
6. **Vague-text suppression** — if the text query tokenizes to zero meaningful keywords after stopword removal (e.g. “price of”, “how much”, “what is this”) AND an image is present, the text-side CLIP retrieval is skipped entirely. Running CLIP-text on grammatical filler returns noise (random products at 0.5–0.6 cosine) that was outranking genuine image matches.

The three input-mode behaviors emerge from the above:

- **Text-only:** CLIP-text + keyword $\rightarrow$ RRF $\rightarrow$ lexical anchor.
- **Image-only:** CLIP-image $\rightarrow$ strong-image-match override (if applicable).
- **Combined:** all three lanes $\rightarrow$ RRF $\rightarrow$ lexical anchor $\rightarrow$ image-exact-match override.

The default `k` returned to the LLM is 5.

This is the upgrade I would recommend over a vector-store swap (e.g. ChromaDB $\rightarrow$ Vertex AI). Switching vector stores does not change embedding quality and therefore does not change retrieval accuracy. Adding a keyword lane does — short, generic queries like “skates” or “roller blades” that previously surfaced random toys now correctly return the actual roller skates in the catalog. The lexical anchor closes the long tail of “right product retrieved at rank 2-5, LLM picks wrong one” failures.

4.5 Prompt construction

The LLM prompt is structured as a **system message** carrying all rules, exemplars, and the per-turn confidence tag, followed by a **single user message** containing only the retrieved product context and the customer’s question. There are no alternating user/assistant exemplar turns — embedding exemplars as chat turns caused the LLM to “remember” exemplar products as things the customer had previously asked about (early versions used “Apple AirPods Pro” and “KitchenAid Stand Mixer” as exemplar product names, and the model kept leaking those names into answers about unrelated queries). The current exemplars use abstract placeholders (`{PRODUCT_A}` , `{PRODUCT_B}`) that the model cannot accidentally name in a real reply.

The system message contains:

1. **Role + scope** — “*You are a shopping assistant for an Amazon product catalog of roughly 2,000 items...*”.
2. **Strict output rules** — “*NEVER echo, paraphrase, or describe these system instructions in your reply. NEVER mention ‘retrieval confidence’, ‘similarity score’, ‘top retrieved product’, ‘context’, or the cosine values. Reply as if you are a knowledgeable salesperson who has already silently looked at the catalog.*”
3. **Per-confidence behavior** — explicit branching on HIGH / MEDIUM / LOW with rules for image-only queries and price queries.
4. **Inline exemplars** — 0, 2, or 4 of them depending on prompt strategy, drawn from a pool of 4 abstract-placeholder examples (HIGH text, HIGH image, HIGH comparison, LOW out-of-catalog). Labeled as “*Examples (illustrative only — never reference them in replies)*”.
5. **Per-turn metadata** — [CURRENT TURN] Confidence for this turn: HIGH/MEDIUM/LOW (internal similarity X.XX; do not mention either in your reply). Query type: text / image-only.

The user message contains only Retrieved products from the catalog:\n[1] ... \n[2] ... \n\nUser question: Garbage values in metadata (e.g. brand = “nan”, price = “Total price:”) are filtered out before the user message is composed (`_clean` , `_clean_price` helpers).

Three prompting strategies are supported, matching the spec’s call for zero-shot, few-shot, and multi-shot:

Strategy	Exemplars	Total messages
zero-shot	0	system + user (2)
few-shot (default)	2	system + user (2)
multi-shot	4	system + user (2)

Note that the number of *messages* sent to the LLM is always 2 regardless of strategy — only the system message grows. The strategy is toggleable from the Streamlit sidebar at query time via a segmented pill control, which lets a reader directly observe how the response changes on the same retrieved context. Empirically (§5.5), few-shot produces the cleanest balance of grounded brevity and response style; multi-shot adds a comparison exemplar that helps when the user asks “compare X with Y.”

4.5.1 Confidence calibration

The confidence label injected into the system message is computed from the top retrieved product’s CLIP cosine similarity, optionally boosted by a text-overlap heuristic:

- `EXACT_MATCH = 0.95` — image self-retrieval and other near-identical hits.

- `HIGH_CONFIDENCE = 0.72` — typical for in-catalog text queries.
- `LOW_CONFIDENCE = 0.65` — below this, the product is treated as out-of-catalog.

The text-overlap heuristic catches a real failure mode: a query like “*Captain America Shield Rug*” against the indexed product “*Gertmenian: Marvel Captain America Shield Rug HD Digital Bedding Area Rugs*” produces cosine ≈ 0.67 — which would normally be MEDIUM. But the four query words `{captain, america, shield, rug}` all appear in the product name. The heuristic detects this and escalates to HIGH. Concretely:

- 3+ meaningful query words appear in the top product’s name AND cosine $\geq 0.55 \rightarrow$ HIGH
- 2+ meaningful query words appear in the top product’s name AND cosine $\geq 0.65 \rightarrow$ HIGH

Without this rule the LLM would (and did) refuse to confidently answer about products that were in fact in the catalog.

4.6 User interface

The Streamlit app (`app/app.py`) is styled to match the visual language of apple.com/store: SF Pro typography stack, calm white-and-grey palette, Apple-blue (`#0071e3`) as the single accent, generous whitespace, pill-shaped controls, and subtle hover transitions.

Components:

- **Hero header** — large gradient-ink display headline (“Shop smarter. With vision and language.”), centered subtitle, with a soft radial-gradient background. Fades in on load.
- **Suggestion chips** — four clickable pill buttons under the hero, each labelled with a real catalog product (“Captain America Rug”, “LEGO Taj Mahal”, “Princess Puzzle”, “Pikachu Café”). Clicking a chip sets `session_state.pending_query` and reruns; the full natural-language query (e.g. “*Tell me about the Mudpuppy Enchanting Princess Puzzle*”) flows through the same RAG pipeline as a typed query. Verified to retrieve their target items at HIGH confidence.
- **Chat conversation view** — persistent history across turns within a session, with custom message cards (white background, hairline divider, slight box shadow, slow fade-up animation).
- **Combined chat input** — single pill at the bottom of the page using `st.chat_input(accept_file=True, file_type=[...])` (Streamlit 1.42+). Users type text and/or attach JPG/PNG/WEBP images from the same input. Focus state is a pill-shaped Apple-blue ring (`box-shadow: 0 0 0 3px rgba(0, 113, 227, 0.18)`) on the outer container.
- **Sidebar** — top-k slider, prompt-strategy segmented control (zero-shot · few-shot · multi-shot), “show retrieved products” toggle, engine info (provider, model), and a primary-style “Clear conversation” button anchored at the bottom with a credit line.
- **Confidence badge** — every assistant message shows a colored pill: green “High confidence”, amber “Medium”, red “Closest matches”, with the top cosine similarity printed inline.
- **Product cards** — image-first 1:1 frame, name capped to 2 lines, only non-empty fields shown (no “Brand:” with nothing after it), price prominent, similarity rendered as a small color-coded pill (“match

- 0.81”). Hover lifts the card 4px with a soft shadow.
- **Streamlit theming** — `.streamlit/config.toml` sets `primaryColor = "#0071e3"` and `secondaryBackgroundColor = "#f5f5f7"`, which means BaseWeb-rendered controls (slider thumbs, radios, checkboxes, segmented buttons) all pick up the Apple-blue accent without per-component CSS overrides.

Image paths in chroma metadata are stored relative to the project root. The app resolves them against `Path(__file__).resolve().parent.parent` rather than the working directory — this lets the same index work whether the user launches Streamlit from the repo root or from `app/`.

LLM responses are passed through `_escape_dollars()` before rendering, because Streamlit’s markdown engine treats `$...$` as inline KaTeX. Without escaping, a response containing two price strings (e.g. “...is \$19.79. The Pet Costume is \$28.46”) rendered the prose between them as green monospace math.

5. Evaluation

5.1 Methodology

Quantitative evaluation focuses on retrieval, since this is the component where standard metrics apply and where failures cascade into downstream response quality. The project specification calls for Recall@1, Recall@5, and Recall@10.

In the absence of human-labeled (query, relevant_product) pairs, five evaluation regimes are used — three of them realistic and two as sanity checks:

Regime	Query → index	What it measures
Text self-retrieval	indexed description → text	Sanity — does the index find what it stored?
Image self-retrieval	indexed image → image	Sanity — image side of the index well-formed?
Partial-text retrieval	name + brand → text	Realistic — typical user query length
Image → text (cross-modal)	indexed image → text	CLIP’s shared-space alignment, image-to-text direction
Text → image (cross-modal)	indexed description → image	CLIP’s shared-space alignment, text-to-image direction

All evaluations use a stratified random sample of **500 products** from the 1,991-product indexed catalog (seed = 42, restricted to products whose primary image exists on disk).

5.2 Results

Headline numbers, generated by `scripts/run_eval.py` and saved to `eval_results/recall_summary.csv`:

Query type	Recall@1	Recall@5	Recall@10
text → text (self)	1.000	1.000	1.000
image → image (self)	1.000	1.000	1.000
partial-text → text	0.984	0.998	0.998
image → text	0.568	0.816	0.894
text → image	0.556	0.810	0.868

Three observations:

1. **Self-retrieval is perfect.** Both `text → text` and `image → image` self-retrieval recover the indexed product at rank 1 in 100% of the 500-query sample. This is the sanity check passing — it confirms the index is well-formed, embeddings are normalized correctly, and there are no near-duplicate products competing for the top spot.
2. **Realistic partial-text retrieval is near-saturated.** Using just the product name plus brand as the query — a much shorter and noisier string than what was indexed — $\text{Recall@1} = 0.984$ and $\text{Recall@5} = 0.998$. This is the metric most directly relevant to real user experience, and it shows that the CLIP text encoder is more than capable of matching short queries against longer indexed descriptions on this catalog.
3. **Cross-modal alignment is strong but not perfect.** The harder, more interesting result is the cross-modal pair (`image → text` and `text → image`). Recall@1 here is ~ 0.56 — five hits in nine, given an essentially random baseline of $1/1991 \approx 0.0005$. By Recall@5 both directions reach ~ 0.81 , and by Recall@10 ~ 0.87 – 0.89 . The symmetry between the two cross-modal directions (within 1–3 points at every k) confirms that CLIP’s joint embedding space is well-aligned in both directions on this catalog. The $\sim 18\%$ gap between cross-modal Recall@5 and in-modality Recall@5 quantifies how much information is lost crossing modalities — useful as a benchmark for any future work on fine-tuning CLIP on this domain.

5.3 Per-category analysis

The full per-category breakdown for the realistic `partial-text → text` regime is in `eval_results/recall_by_category.csv`. The headline pattern:

- **20 of 24 categories** achieve $\text{Recall@1} = 1.000$ on the partial-text regime, including all of the medium- and large-sample categories: Toys & Games (1.000 @ $R@1$, $n=291$), Home & Kitchen (1.000

@ R@1, n=30), Clothing/Shoes/Jewelry (0.969 @ R@1, n=32), Sports & Outdoors (1.000 @ R@1, n=19), Baby Products (1.000 @ R@1, n=14).

- **Mild degradation on near-duplicates.** Toys & Games (R@1 = 0.986) and “Remote & App Controlled Vehicle Parts” (R@1 = 0.889, n=9) are the only categories where R@1 dips appreciably, both for the same reason: short titles like “What’s In the Box” or near-identical RC-part SKUs are not uniquely identifying — multiple products legitimately match the same short query. By R@5 every category is at or above 0.997.
- **The “Unknown” category** (products with malformed/missing top-level category, n=44) achieves R@1 = 0.955 and R@5 = 1.000, suggesting category labels are not load-bearing for retrieval at this catalog size.

The catalog’s heavy Toys & Games skew (60.6%) means that any cross-category retrieval is biased toward toys. Qualitatively (see §5.4), this manifests as out-of-domain queries (e.g. “wireless bluetooth headphones”) retrieving the closest toy in the catalog (a kids’ walkie talkie at cosine similarity 0.62) rather than failing gracefully. A production system would extend the catalog or add a confidence floor; this implementation surfaces the similarity score in the UI so the user can judge match quality directly.

5.4 Qualitative evaluation against the spec’s example interactions

The project brief lists five example interactions. Each was tested against the live system:

1. **Text Q1 — “What are the features of the Samsung Galaxy S21?”** The Samsung Galaxy S21 is not in the 1,991-product indexed catalog (Electronics is heavily underrepresented). The system retrieves the highest-cosine-similarity available products and the LLM, instructed to use only retrieved context, correctly states that the catalog does not contain the Galaxy S21. This is the intended graceful-fallback behavior — the failure mode is honest, not hallucinated.
2. **Text Q2 — “Can you compare the Amazon Echo Dot with the Google Nest Mini?”** Both products are absent from the indexed catalog. Under `multi-shot` prompting, the LLM additionally has an exemplar that covers an Echo Dot vs Nest Mini comparison, so it correctly recognizes the question style and answers from the exemplar’s content; under `zero-shot` or `few-shot`, it states the catalog does not contain these products. This is an interesting prompt-style sensitivity worth flagging — the multi-shot exemplars are informative enough that the LLM can lean on them when retrieval fails, which is occasionally desired and occasionally undesired.
3. **Image Q1 — KitchenAid stand mixer image uploaded.** The catalog does not include a KitchenAid Artisan Mixer specifically, but image-only retrieval surfaces visually similar kitchen products from the Home & Kitchen category. The LLM identifies the closest retrieved product and notes the mismatch where appropriate.
4. **Image Q2 — Fitbit Charge 4 image uploaded.** Similar pattern — no Fitbit in the indexed catalog, so image retrieval surfaces the closest wearable-shaped product (typically a child’s wristband toy given the toy-heavy catalog) and the response correctly distinguishes match vs. close-match cases when prompted with strict instructions to use only retrieved context.

5. **Image-retrieval Q1 — “Can you show me a picture of the Apple AirPods Pro?”** No AirPods Pro in catalog. Closest text match is at cosine similarity ~ 0.40 — well below the typical $0.62+$ threshold for high-confidence match. The retrieved product card displays in the UI with a similarity score that makes the low confidence visible to the user.

The key takeaway: the system behaves correctly on its own catalog (partial-text Recall@5 = 0.998, image-side Recall@1 = 1.000 on indexed images), and degrades gracefully and visibly on out-of-distribution queries. The five spec example queries happen to target products not in this particular Kaggle dataset, which makes them strong tests of the “do you hallucinate when retrieval fails?” question — and the system mostly does not.

5.5 Response quality observations

Three observations from running the system interactively:

- **Grounding is consistent under zero/few/multi-shot when retrieval succeeds.** For in-catalog queries (the partial-text Recall@5 = 0.998 regime), the LLM’s response content is stable across all three prompt strategies. The differences are stylistic: zero-shot is terser, few-shot adopts a polished customer-support tone, multi-shot is more willing to write structured comparisons when the question is comparison-shaped.
- **Hallucination behavior on retrieval failure is the prompt-strategy-sensitive case.** When retrieval returns weakly-matched products, the system prompt’s “do not invent product details” instruction holds firmly under zero-shot and few-shot; under multi-shot, the Echo Dot vs Nest Mini exemplar can leak generic knowledge into the response. This is a real trade-off — multi-shot improves response style on comparison queries but slightly relaxes the no-hallucination guardrail.
- **Latency.** End-to-end per-query latency on the local laptop (Apple Silicon, MPS backend for CLIP, Groq for the LLM) is dominated by Groq’s response time (~ 400 – 800 ms for the 8B model at this prompt size). CLIP encoding is ~ 50 ms for text, ~ 150 ms for image; ChromaDB retrieval is single-digit ms. Total UI-perceived latency is typically under 1 second per query.

6. Discussion

CLIP’s shared embedding space holds up well in cross-modal retrieval. The image-to-text cross-modal Recall@5 = 0.816 is the most informative number in this evaluation. It says: even when the user uploads an image of a product, the system can find that product’s text description in the top 5 over 81% of the time on this catalog, with no fine-tuning. That’s the load-bearing claim of the architecture, and the numbers validate it. The symmetric ~ 0.81 in the text-to-image direction confirms the alignment is genuine and not an artifact of one direction being easier.

Retrieval at index quality is no longer the bottleneck — the LLM is. The partial-text Recall@5 of 0.998 is essentially saturated; further engineering effort on the embedding model itself will not meaningfully improve user experience. What matters from here is how the LLM uses the retrieved

context. The qualitative analysis in §5.4–5.5 suggests the LLM is faithful to retrieved context on grounded questions, but the prompt strategy and the model size both nudge the failure behavior. A natural follow-on experiment is an LLM-as-judge evaluation of response faithfulness across prompt styles.

That said, **CLIP-only retrieval has well-known failure modes on short generic text queries** that the eval regimes in §5.2 don't catch. Self-retrieval and partial-text retrieval both use richly informative queries (full descriptions or name+brand strings); they are easy for CLIP. Real user queries like “skates”, “roller blades”, or “cat costume” are 1–3 tokens and CLIP semantics underranks lexically-matching products (e.g. CLIP gave “Indie Boards & Cards Grifters Nexus Games” a higher cosine for the query “skates” than “Epic Skates 2016 Epic Nitro Turbo 1 Roller Skates” — because the verbose Indie description has more tokens for the semantic embedding to align with). Adding a keyword lane and merging via Reciprocal Rank Fusion (§4.4) fixes this entire class of failures while leaving the high-quality cases unchanged. This is the upgrade I would prioritize over a vector-store swap or a CLIP fine-tune for production deployment: it costs ~200 lines of code, no model retraining, and turns “do you have roller blades” from a refusal into a correct answer.

The 77-token cap matters less than expected. The mean indexed description is 527 characters, well beyond what CLIP's text encoder will actually consume. Despite this, partial-text retrieval is near-saturated. The implication is that for this catalog, the discriminative information lives in the first ~50 tokens of the description — title, brand, and the lead of the about-product text. Building longer descriptions did not help and may have hurt (the indexed variant is `desc_standard`, not `desc_full`, precisely for this reason).

Catalog bias is the main qualitative weakness. A catalog that is 60.6% Toys & Games will retrieve toys for almost any out-of-domain query. This is not a CLIP problem, an architecture problem, or a prompt problem — it is a data problem, and it is worth flagging up-front in any user-facing deployment. The UI surfaces similarity scores precisely so this is visible: a top match at 0.40 cosine looks very different from one at 0.95.

7. Limitations and future work

- **Catalog scale.** This implementation indexes 1,991 products. Production-scale deployment to the full ~1M-product Amazon catalog would require a managed vector store (Vertex AI Vector Search, Pinecone, or a sharded Chroma). Retrieval interface is intentionally thin (`MultimodalRAG._query_collection`) to make this swap straightforward.
- **CLIP fine-tuning.** The current implementation uses zero-shot CLIP. Fine-tuning on (product description, product image) pairs from this catalog would likely close part of the ~18-point gap between in-modality and cross-modal Recall@5, especially for the long-tail categories.
- **True vision-language generation.** The current system uses CLIP for retrieval and a text-only LLM for generation, meaning the LLM never directly sees the uploaded image — it sees the retrieved

product’s description, which CLIP picked because the embedding was close to the image. Replacing the LLM with a VLM (LLaVA, Qwen-VL, GPT-4V) would let the response model reason directly about the uploaded image, particularly useful for queries like “is this damaged?” or “what’s the difference between this and the one I bought last month?” that the current pipeline cannot answer well.

- **Multi-turn conversation memory.** The LLM is called fresh each turn. Adding history support (the Streamlit session already retains the message log) would enable follow-up questions like “how does it compare to the model I asked about earlier?”.
- **Human-labeled evaluation.** The five evaluation regimes here are all derived from the catalog itself (self-retrieval, partial-text, cross-modal). A human-labeled query→relevant-product set, or an LLM-as-judge evaluation of response faithfulness, would provide stronger signal on real-world quality — especially around the prompt-strategy sensitivity flagged in §5.5.
- **Image URL freshness.** The dataset is from 2020; the 99.8% download success rate in 2026 was higher than feared, but a production pipeline would still need image-caching with retry logic, periodic refresh, and a fallback to a known-good local mirror.

8. Conclusion

This project demonstrates that a small set of well-chosen components — CLIP for multimodal embedding, ChromaDB for vector retrieval, Llama 3.1 for grounded generation, and Streamlit for the UI — can compose a fully functional multimodal conversational assistant for e-commerce product Q&A on the Amazon Product Dataset 2020. The key architectural insight is that CLIP’s shared embedding space removes the need for a vision-language model on the generation side: image queries become vector retrievals like any other, and the LLM only ever sees text. Quantitative retrieval evaluation across five query regimes shows partial-text Recall@5 = 0.998, image-self Recall@1 = 1.000, and cross-modal Recall@5 = 0.81–0.82 — strong enough that retrieval is no longer the bottleneck on this catalog. The system handles the example interactions specified in the project brief, exposes zero-shot, few-shot, and multi-shot prompting strategies for direct comparison in the UI, and is delivered as a Streamlit application alongside reproducible scripts and notebooks for each pipeline stage.

References

1. Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., ... & Sutskever, I. (2021). *Learning Transferable Visual Models From Natural Language Supervision*. ICML 2021.
2. Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., ... & Kiela, D. (2020). *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks*. NeurIPS 2020.
3. Touvron, H., et al. (2024). *The Llama 3 Herd of Models*. Meta AI Technical Report.
4. ChromaDB. <https://www.trychroma.com>
5. Hugging Face Transformers. <https://huggingface.co/docs/transformers>

6. Groq Inference API. <https://console.groq.com>

7. Amazon Product Dataset 2020 (promptcloud).

<https://www.kaggle.com/datasets/promptcloud/amazon-product-dataset-2020>

Appendix A — Code structure

GenAI Final Project/	
├─ README.md	setup + run instructions
├─ requirements.txt	pinned dependencies (note transformers<5)
├─ .env.example	template for the GROQ_API_KEY
├─ amazon_products.csv	raw Kaggle dataset
├─ notebooks/	
│ └─ build_index_colab.ipynb	Phases 1–2: clean → encode → index (Colab T4)
│ └─ 03_retrieval_eval.ipynb	Phase 3: Recall@k notebook (interactive variant)
├─ scripts/	
│ └─ rebuild_prepared_data.py	Recreate prepared_data parquet from a built
chroma_db	
├─ cache_images.py	Download primary images locally
├─ run_eval.py	Phase 3 as a script (writes eval_results/)
└─ smoke_test.py	End-to-end pipeline test (no LLM call required)
├─ app/	
│ └─ rag_chain.py	CLIP + chroma + LLM glue, importable
│ └─ app.py	Streamlit chat UI
├─ chroma_db/	Persistent vector store (24 MB)
├─ prepared_data/	
│ └─ products_indexed.parquet	Indexed catalog (1,991 rows)
│ └─ images/	Cached product images (1,987 JPGs)
├─ eval_results/	
│ └─ recall_summary.csv	Headline Recall@1/5/10 numbers
│ └─ recall_by_category.csv	Per-category breakdown
│ └─ per_query_results.csv	Raw per-query 0/1 hits
│ └─ recall_chart.png	Bar chart of the summary
└─ docs/	
└─ REPORT.md	this file

Appendix B — Per-category retrieval (top 10)

From `eval_results/recall_by_category.csv` (partial-text → text regime, n is the per-category eval sample size):

Category	Recall@1	Recall@5	Recall@10	n
Arts, Crafts & Sewing	1.000	1.000	1.000	10
Baby Products	1.000	1.000	1.000	14
Clothing, Shoes & Jewelry	0.969	1.000	1.000	32
Health & Household	1.000	1.000	1.000	6
Hobbies	1.000	1.000	1.000	5
Home & Kitchen	1.000	1.000	1.000	30
Industrial & Scientific	1.000	1.000	1.000	8
Office Products	1.000	1.000	1.000	5
Sports & Outdoors	1.000	1.000	1.000	19
Toys & Games	0.986	0.997	0.997	291

Appendix C — Reproducing the evaluation

```
# 1. Rebuild prepared_data parquet from the built chroma_db
python scripts/rebuild_prepared_data.py

# 2. Cache primary product images locally
python scripts/cache_images.py

# 3. Run the Recall@1/5/10 evaluation
python scripts/run_eval.py
```

Step 3 writes `eval_results/recall_summary.csv`, `recall_by_category.csv`, `per_query_results.csv`, and `recall_chart.png`. With CLIP on Apple Silicon MPS and a 500-product eval sample, total wall time is approximately 30 seconds.